




Clinical-Trial Eligibility Copilot

Human-reviewed AI support for criterion-level clinical-trial screening

Eligibility screening connects patient recruitment with enrolment



AI-supported patient-trial matching

Criterion-level eligibility review with human-review routing

👁 Target client

Mid-sized clinical research organisations, sponsors, and CROs managing multiple active studies with recruitment teams and complex patient information

👁 Market signal

Deep 6 AI, Tempus, TrialX, and WCG: AI-supported patient-trial matching or pre-screening across structured and unstructured clinical data.

👁 Use cases considered

- Patient-trial discovery
- Recruitment prioritisation
- Criterion-level eligibility review

AI-supported matching is already a market reality

Faster screening with visible human control

The product opportunity - Trial eligibility prescreening - Safe & Efficient

👁️ **Current process**

Clinical coordinators manually compare narrative patient information with protocol criteria. The work is repetitive, difficult to scale, and hard to document consistently.

Value driver	Proposed pilot metric	Why an operations lead cares
Coverage	Potential candidates assessed	More eligible candidates reviewed across active studies
Speed	Time from identification to coordinator review	Recruitment teams can focus effort where necessary
Safety	Uncertain and high-risk cases escalated	Ambiguous or potentially unsafe outputs are not treated silently
Trust	Assessments with rationale, evidence, and trace	Coordinators understand, challenge, and document the AI-supported assessment

👁️ **Value hypothesis**

Reduce screening burden while preserving traceability, escalation of uncertainty, and mandatory human clinical review

Performance, safety, workload, and cost

Early evidence

AI-assisted clinical-trial pre-screening | Evaluation dashboard


120 public synthetic TrialGPT assessments · GPT-4.1 · Decision support only — human review remains mandatory

Exact agreement 72,5%	Unsafe MET rate 5,9%	UNKNOWN recall 92,3%	Review rate 59,2%	Median latency 1,01s	Cost per assessment \$0,0021
--------------------------	-------------------------	-------------------------	----------------------	-------------------------	---------------------------------


LLM prediction vs reference label

Ground Truth Label	Predicted Label			
	MET	NOT_APPLICAB..	NOT_MET	UNKNOWN
MET	24		1	7
NOT_APPLICABLE	3	3	4	9
NOT_MET	1		12	4
UNKNOWN	1		3	48


LLM agreement by criterion type



LLM review workload by criterion type



Reference outcome distribution - 120 assessments



Evaluation cohort: 120 criterion-level assessments across 30 public synthetic TrialGPT patient summaries.

Model: GPT-4.1 · Prompt: v2_abstention_rules · Estimated API cost: \$0.0021 per assessment.

Source: NCBI TrialGPT Criterion Annotations. Decision-support prototype only; final eligibility and enrolment decisions require human clinical-trial coordinator review.

The opportunity is not autonomous screening; it is faster, evidence-supported preparation for human review

Human review by design


POC

Synthetic TrialGPT input

Python + GPT4.1
Criterion assessment

N8n webhook routing logic

LangSmith monitoring



	Name	Input	Output	Error	Start Time	Latency	Dataset	Annotation ...	Tokens	Cost	First Token
	ChatOpenAI	{ "criterion_text": "DSM-I...	message: ""		8/30/2026, 1:3...	1.02s			773	\$0.001918	
	criterion_level_screening	gpt-4.1	[{"criterion_id": "CRIT_063...		8/30/2026, 1:3...	1.02s			773	\$0.001918	
	ChatOpenAI	You are an AI assistant s...	resp_015c77d29c512330...		8/30/2026, 1:3...	1.28s			759	\$0.00186	
	criterion_level_screening	gpt-4.1	[{"criterion_id": "CRIT_069...		8/30/2026, 1:3...	1.29s			759	\$0.00186	
	ChatOpenAI	{ "criterion_text": "are not ...	message: ""		8/30/2026, 1:3...	2.08s			865	\$0.0023	
	criterion_level_screening	gpt-4.1	[{"criterion_id": "CRIT_042...		8/30/2026, 1:3...	2.16s			865	\$0.0023	

Notion review queue

Clinical trial eligibility Copilot						
Clinical-trial eligibility review queu...						
Review case	Queue status	Route	Patient ID	Trial ID	Queue ID	Criterion ID
sigir-201425 NCT00282269	OPEN	HUMAN_REVIEW	sigir-201425	NCT00282269	poc-unknown-002	CRIT_042
sigir-201425 NCT00282269	OPEN	HUMAN_REVIEW	sigir-201425	NCT00282269	poc-unknown-001	CRIT_042
sigir-201425 NCT00282269	OPEN	HUMAN_REVIEW	sigir-201425	NCT00282269	poc-unknown-001	CRIT_042

An 8-week controlled pilot can test value before production investment

Indicative cost, timeline, assumptions and next-step ask

- ▶ Indicative upfront cost
€23,408
Controlled pilot covering workflow configuration, evaluation, monitoring, documentation and project coordination.

Timing	Activity	Output
Week 1	Scope and clinical-rule review	Agreed workflow, escalation rules and success criteria
Weeks 2–3	Data mapping and governance	Approved pilot data and controlled input process
Weeks 3–4	Workflow and monitoring configuration	Review routing, audit trail and dashboard
Weeks 5–6	Human-reviewed pilot	Safety, workload and time-saving evidence
Weeks 7–8	Evaluation and decision	Continue, pivot or stop recommendation

Key assumptions


- One bounded use case using synthetic or approved pilot data.
- Human review remains mandatory for every clinical decision.
- Existing POC tools are used; enterprise licences and full EHR/CTMS integration are excluded.

API usage is inexpensive; clinical validation, governance, integration and human-review effort drive the real cost

The goal is not to replace clinical judgement,
but to help experts reach safer, faster and
more transparent decisions

Human review by design

Production data flow



Eligibility-copilot POC

Simulated upstream data

Public synthetic TrialGPT case
Structured patient, trial and
criterion data

Criterion-level AI assessment



Guardrails and routing



Monitoring across the flow

